



```
In [ ]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [ ]: df = pd.read_csv("/content/Exam_Score_Prediction.csv")  
df.head()
```

```
Out[ ]:   student_id  age  gender  course  study_hours  class_attendance  internet_access  
0           1    17     male  diploma        2.78            92.9  
1           2    23    other    bca        3.37            64.8  
2           3    22     male    b.sc        7.88            76.8  
3           4    20    other  diploma        0.67            48.4  
4           5    20   female  diploma        0.89            71.6
```

```
In [ ]: df.shape
```

```
Out[ ]: (20000, 13)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20000 entries, 0 to 19999  
Data columns (total 13 columns):  
 #   Column          Non-Null Count  Dtype     
---  --  
 0   student_id      20000 non-null  int64    
 1   age             20000 non-null  int64    
 2   gender          20000 non-null  object    
 3   course          20000 non-null  object    
 4   study_hours     20000 non-null  float64  
 5   class_attendance 20000 non-null  float64  
 6   internet_access 20000 non-null  object    
 7   sleep_hours     20000 non-null  float64  
 8   sleep_quality   20000 non-null  object    
 9   study_method    20000 non-null  object    
 10  facility_rating 20000 non-null  object    
 11  exam_difficulty 20000 non-null  object    
 12  exam_score      20000 non-null  float64  
dtypes: float64(4), int64(2), object(7)  
memory usage: 2.0+ MB
```

```
In [ ]: df.describe()
```

Out[]:

	student_id	age	study_hours	class_attendance	sleep_hours
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	10000.504600	20.473300	4.007604	70.017365	7.00856
std	5773.654959	2.284458	2.308313	17.282262	1.73209
min	1.000000	17.000000	0.080000	40.600000	4.10000
25%	5000.750000	18.000000	2.000000	55.100000	5.50000
50%	10000.500000	20.000000	4.040000	69.900000	7.00000
75%	15000.250000	22.000000	6.000000	85.000000	8.50000
max	20001.000000	24.000000	7.910000	99.400000	9.90000

In []: `df.isnull().sum()`

Out[]:

	0
student_id	0
age	0
gender	0
course	0
study_hours	0
class_attendance	0
internet_access	0
sleep_hours	0
sleep_quality	0
study_method	0
facility_rating	0
exam_difficulty	0
exam_score	0

dtype: int64

In []: `df.columns`

Out[]: `Index(['student_id', 'age', 'gender', 'course', 'study_hours', 'class_attendance', 'internet_access', 'sleep_hours', 'sleep_quality', 'study_method', 'facility_rating', 'exam_difficulty', 'exam_score'], dtype='object')`

This dataset has no null values so our data is cleaned

```
In [ ]: df.dtypes
```

```
Out[ ]:          0
student_id    int64
age           int64
gender        object
course         object
study_hours   float64
class_attendance  float64
internet_access  object
sleep_hours   float64
sleep_quality  object
study_method   object
facility_rating  object
exam_difficulty  object
exam_score    float64
```

dtype: object

```
In [ ]: df = df.drop("student_id", axis=1)
```

```
In [ ]: df.columns
```

```
Out[ ]: Index(['age', 'gender', 'course', 'study_hours', 'class_attendance',
               'internet_access', 'sleep_hours', 'sleep_quality', 'study_method',
               'facility_rating', 'exam_difficulty', 'exam_score'],
              dtype='object')
```


Label encoding

```
In [ ]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
categorical_cols = [
    "gender", "course", "internet_access", "sleep_quality",
    "study_method", "facility_rating", "exam_difficulty"
]


for col in categorical_cols:
    df[col]=le.fit_transform(df[col])
```

Data Visualization


```
In [ ]: sns.histplot(df["exam_score"], kde=True)
plt.title("Distribution of Exam Scores")
plt.show()
```



```
In [ ]: plt.figure(figsize=(8,6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [ ]: sns.scatterplot(x="study_hours", y="exam_score", data=df)
plt.title("Study Hours vs Exam Score")
plt.show()
```



```
In [ ]: df["Result"] = df["exam_score"].apply(
    lambda x: 1 if x >= 50 else 0
)
```

```
In [ ]: df = df.drop("exam_score", axis=1)
```

SVM

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [ ]: #Extracting datasets
x = df.iloc[:, :-1].values #independent
y = df.iloc[:, -1].values #dependent
```

```
In [ ]: #Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random
```

```
In [ ]: # Feature Scaling
from sklearn.preprocessing import StandardScaler
st_x=StandardScaler()
x_train=st_x.fit_transform(x_train)
```

```
x_test=st_x.transform(x_test)
```

```
In [ ]: from sklearn.svm import SVC #SUPPORT VECTOR CLASSIFIER
classifier= SVC(kernel='linear',random_state=0)
classifier.fit(x_train,y_train)
```

```
Out[ ]: SVC
SVC(kernel='linear', random_state=0)
```

```
In [ ]: #Predicting the test set result
y_pred= classifier.predict(x_test)
```

```
In [ ]: #Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test, y_pred)
cm
```

```
Out[ ]: array([[ 707,  398],
   [ 256, 2639]])
```

```
In [ ]: from sklearn.metrics import accuracy_score, classification_report
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.8365

Classification Report:					
	precision	recall	f1-score	support	
0	0.73	0.64	0.68	1105	
1	0.87	0.91	0.89	2895	
accuracy			0.84	4000	
macro avg	0.80	0.78	0.79	4000	
weighted avg	0.83	0.84	0.83	4000	

Now SVM ON PARTICULAR COLUMNS

```
In [ ]: selected_features = [
    "study_hours",
    "class_attendance",
    "study_method",
    "sleep_hours",
    "sleep_quality",
    "internet_access",
    "exam_difficulty",
    "age",
    "gender"
]
```

```
In [ ]: X = df[selected_features].values  
y = df["Result"].values
```

```
In [ ]: from sklearn.model_selection import train_test_split  
  
x_train, x_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=0  
)
```

```
In [ ]: from sklearn.preprocessing import StandardScaler  
  
sc = StandardScaler()  
x_train = sc.fit_transform(x_train)  
x_test = sc.transform(x_test)
```

```
In [ ]: from sklearn.svm import SVC  
  
classifier = SVC(kernel='linear', random_state=0)  
classifier.fit(x_train, y_train)
```

```
Out[ ]: ▾ SVC  
SVC(kernel='linear', random_state=0)
```

```
In [ ]: #Predicting the test set result  
y_pred = classifier.predict(x_test)
```

```
In [ ]: #Creating the Confusion matrix  
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
cm
```

```
Out[ ]: array([[ 709,  396],  
               [ 254, 2641]])
```

```
In [ ]: from sklearn.metrics import accuracy_score, classification_report  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.8375

Classification Report:				
	precision	recall	f1-score	support
0	0.74	0.64	0.69	1105
1	0.87	0.91	0.89	2895
accuracy			0.84	4000
macro avg	0.80	0.78	0.79	4000
weighted avg	0.83	0.84	0.83	4000

```
In [ ]: selected_features = [
    "study_hours",
    "class_attendance",
    "study_method",
    "sleep_hours",
    "sleep_quality",
    "internet_access",
    "exam_difficulty",
    "age",
    "gender"
]
```

```
In [ ]: X = df[selected_features].values
y = df["Result"].values
```

```
In [ ]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=0
)
```

```
In [ ]: from sklearn.preprocessing import StandardScaler

sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
In [ ]: from sklearn.svm import SVC

classifier = SVC(kernel='rbf', C=1, gamma='scale', random_state=0)
classifier.fit(x_train, y_train)
```

```
Out[ ]: ▾ SVC ⓘ ?
```

SVC(C=1, random_state=0)

```
In [ ]: #Predicting the test set result
y_pred = classifier.predict(x_test)
```

```
In [ ]: #Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

```
Out[ ]: array([[ 699,  406],
               [ 226, 2669]])
```

```
In [ ]: from sklearn.metrics import accuracy_score, classification_report
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.842

Classification Report:					
	precision	recall	f1-score	support	
0	0.76	0.63	0.69	1105	
1	0.87	0.92	0.89	2895	
accuracy			0.84	4000	
macro avg	0.81	0.78	0.79	4000	
weighted avg	0.84	0.84	0.84	4000	

LOGISTIC REGRESSION

TAKING ALL COLUMNS

```
In [ ]: #Extracting datasets
x = df.iloc[:, :-1].values #independent
y = df.iloc[:, -1].values #dependent
```

```
In [ ]: # Splitting the dataset into training and test set.
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.25, rand
```

```
In [ ]: #feature Scaling
#column=> mean=>
from sklearn.preprocessing import StandardScaler
st_x= StandardScaler()
x_train= st_x.fit_transform(x_train)
x_test= st_x.transform(x_test)
```

```
In [ ]: #Fitting Logistic Regression to the training set
from sklearn.linear_model import LogisticRegression
classifier= LogisticRegression(random_state=0)
classifier.fit(x_train, y_train)
```

```
Out[ ]: ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression(random_state=0)
```

```
In [ ]: #Predicting the test set result
y_pred= classifier.predict(x_test)
```

```
In [ ]: #Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_pred)
cm
```

```
Out[ ]: array([[ 884,  490],  
   [ 326, 3300]])
```

```
In [ ]: from sklearn.metrics import accuracy_score  
  
accuracy = accuracy_score(y_test, y_pred)  
print(f"Accuracy: {accuracy:.2f}")
```

```
Accuracy: 0.84
```

TAKING PARTICULAR COLUMNS

```
In [ ]: selected_features = [  
    "study_hours",  
    "class_attendance",  
    "study_method",  
    "sleep_hours",  
    "sleep_quality",  
    "internet_access",  
    "exam_difficulty",  
    "age",  
    "gender"  
]
```

```
In [ ]: X = df[selected_features].values  
y = df["Result"].values
```

```
In [ ]: # Splitting the dataset into training and test set.  
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.25, rand
```

```
In [ ]: #feature Scaling  
#column=> mean=>  
from sklearn.preprocessing import StandardScaler  
st_x= StandardScaler()  
x_train= st_x.fit_transform(x_train)  
x_test= st_x.transform(x_test)
```

```
In [ ]: #Fitting Logistic Regression to the training set  
from sklearn.linear_model import LogisticRegression  
classifier= LogisticRegression(random_state=0)  
classifier.fit(x_train, y_train)
```

```
Out[ ]: ▾ LogisticRegression ⓘ ?  
LogisticRegression(random_state=0)
```

```
In [ ]: #Predicting the test set result  
y_pred= classifier.predict(x_test)
```

```
In [ ]: #Creating the Confusion matrix
```

```
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_pred)
cm
```

```
Out[ ]: array([[ 884,  490],
 [ 326, 3300]])
```

```
In [ ]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

```
Accuracy: 0.84
```

```
In [ ]: corr_matrix = df.corr()
```

```
In [ ]: corr_with_target = corr_matrix["Result"].sort_values(ascending=False)
print(corr_with_target)
```

```
Result          1.000000
study_hours     0.561711
class_attendance 0.221395
sleep_hours      0.096430
exam_difficulty  0.001381
age              0.000972
course           -0.000088
gender            -0.001972
internet_access   -0.002469
facility_rating    -0.046881
sleep_quality     -0.081065
study_method       -0.111482
Name: Result, dtype: float64
```

```
In [ ]: selected_features_corr = [
        "study_hours",
        "class_attendance"
    ]
```

```
In [ ]: X = df[selected_features].values
y = df["Result"].values
```

```
In [ ]: # Splitting the dataset into training and test set.
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.25, rand
```

```
In [ ]: #feature Scaling
#column=> mean=>
from sklearn.preprocessing import StandardScaler
st_x= StandardScaler()
x_train= st_x.fit_transform(x_train)
x_test= st_x.transform(x_test)
```

```
In [ ]: #Fitting Logistic Regression to the training set
```

```
from sklearn.linear_model import LogisticRegression
classifier= LogisticRegression(random_state=0)
classifier.fit(x_train, y_train)
```

Out[]:

```
▼ LogisticRegression ⓘ ?  
LogisticRegression(random_state=0)
```

In []:

```
#Predicting the test set result
y_pred= classifier.predict(x_test)
```

In []:

```
#Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_pred)
cm
```

Out[]:

```
array([[ 884,  490],
       [ 326, 3300]])
```

In []:

```
from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

Accuracy: 0.84